



# ADVANCE DATABASE SYSTEMS DEVELOPMEN T [CC6001] WEEK - 04

## Database Recovery II

# Definition

---

Database Recovery is the process  
of **RESTORING** the database to the  
**MOST RECENT CONSISTENT STATE**  
that **EXISTED**  
Just before the **FAILURE**

# Definition

---

“restore database to a stable  
that is **known** to be correct  
after some failure has rendered  
current state incorrect or at  
least **suspect**”

# Types of Failures Contd..

---

- ❑ Transaction Failure
- ❑ Media Failure
- ❑ System Failure

# Recovery – System Failure

---

Will affect ALL TRANSACTIONS currently in PROGRESS but NOT database itself.

System recovery is carried out as part of system's restart procedure.

# Techniques for Recovery

---



Deferred Update

ImmediateUpdate

Shadow Paging

# Key Point about System Failure

---

System Failure results in loss  
of content of Main Memory  
[Database Buffer]

# Key Points

Any transaction in progress at the time of failure can not be successfully completed

So, these transactions must be **UNDONE** [rollback] when system starts





# Key Points

Any transaction which successfully completed BUT were not COMMITTED

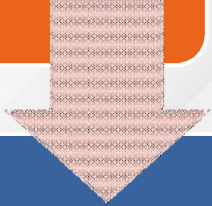
So, these transactions must be REDO when system starts



# Checkpoint

---

Log Files may also contain  
CHECKPOINT records



CHECKPOINT is a certain  
prescribed or scheduled time  
interval system that takes

# At save point

---

Physically  
writes the  
content of  
database buffer  
to physical  
database

Physically writes a  
special checkpoint  
record to the log

Checkpoint record  
gives a list of  
transactions in  
progress at the  
time the savepoint  
is taken

# Deferred Update

Under this recovery protocol, UPDATES are not WRITTEN to database until TRANSACTION has reached its COMMIT point

If TRANSACTION fails before reaching its COMMIT point, it doesn't modify database, hence NO action is needed.

If SYSTEM FAILURE occurs after COMMIT point, REDO the updates of COMMITTED TRANSACTION

# Intermediate Update

---

- Under this recovery protocol, UPDATES are WRITTEN to the database immediately without waiting to reach the COMMIT point
- If TRANSACTION fails, the system should UNDO the updates made by TRANSACTION not COMMITTED at the time of failure.
- May be necessary to REDO the updates of COMMITTED TRANSACTIONS

# Intermediate Update – Key Operations

## REDO

- Transaction which completed successfully before the crash/failure

## UNDO

- Transaction which started but did not complete before the crash/failure

# Intermediate Update – Process

## UNDO List

- If begin transaction is found
- Backward Direction

## REDO List

- If commit is found, move from UNDO to REDO
- Forward Direction

# Intermediate Update - Example

## Log

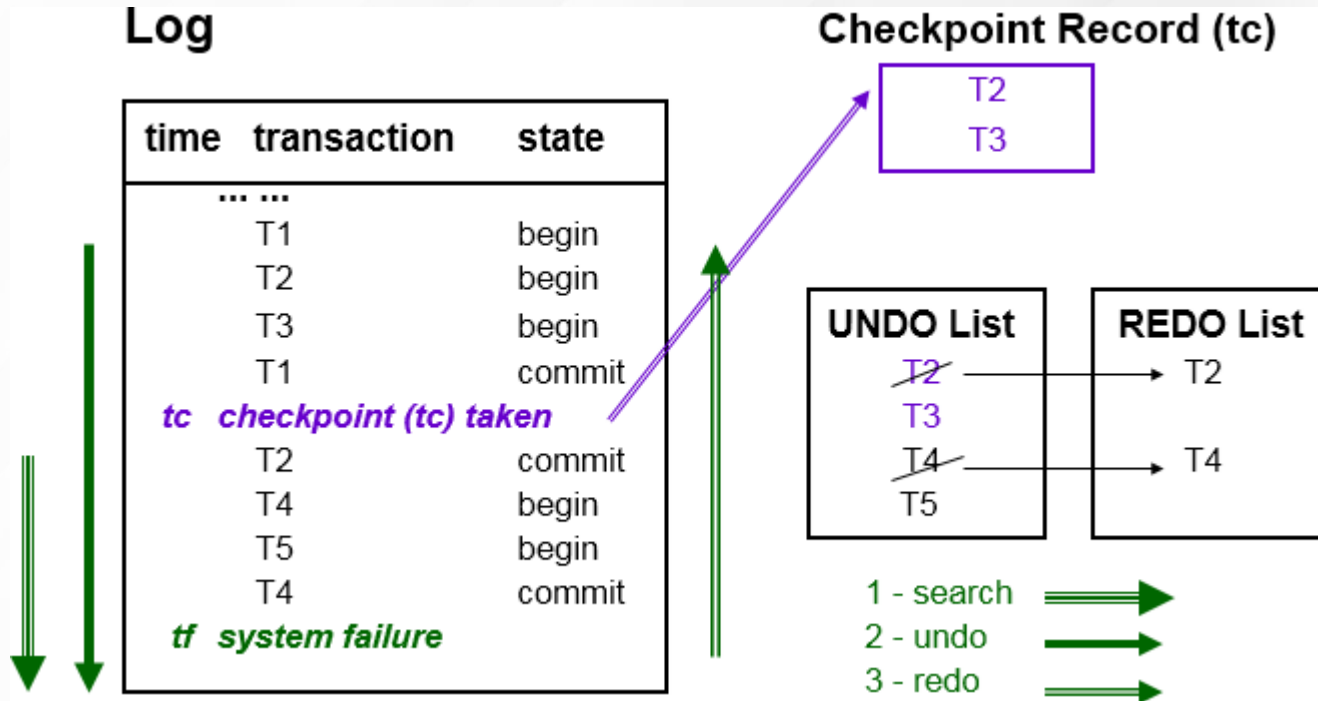
time	transaction	state
...	...	...
	T1	begin
	T2	begin
	T3	begin
	T1	commit
<i>tc</i>	<i>checkpoint (tc) taken</i>	
	T2	commit
	T4	begin
	T5	begin
	T4	commit
<i>tf</i>	<i>system failure</i>	

## Checkpoint Record (tc)

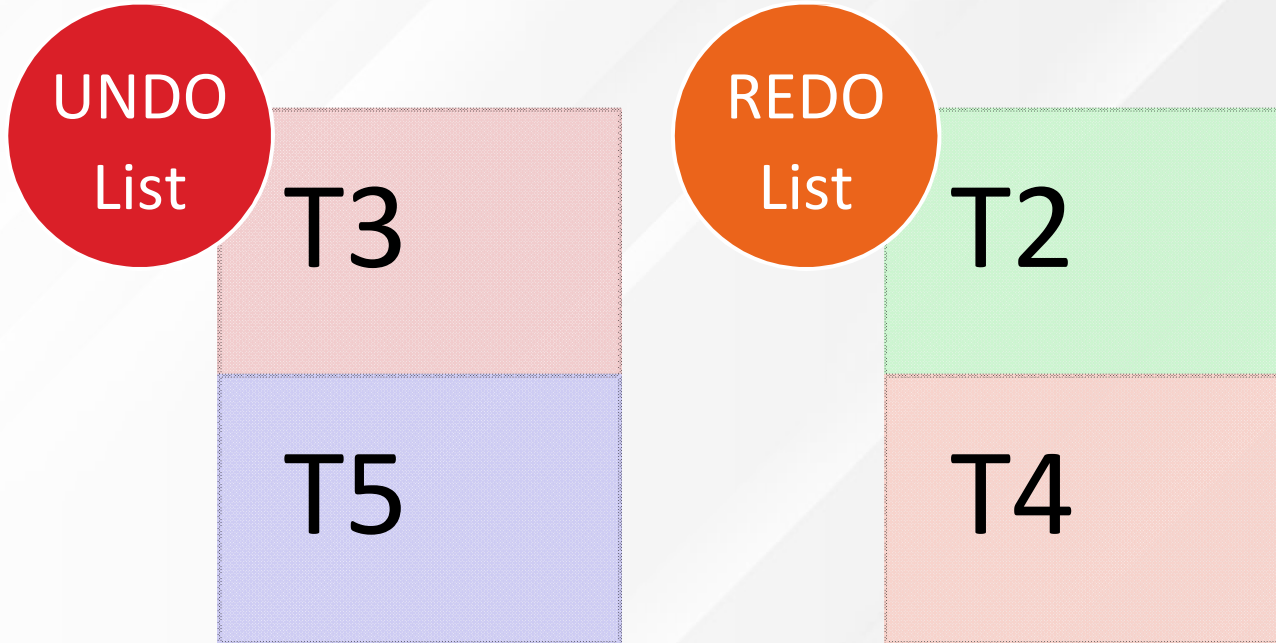
T2  
T3



# Intermediate Update - Example

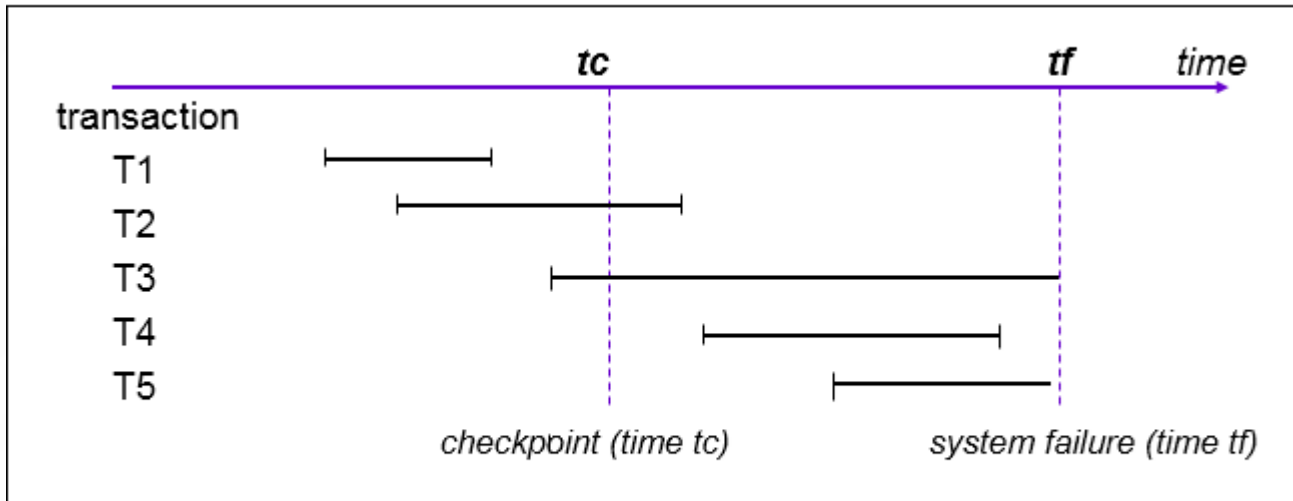


# Intermediate Update - Solution



# Intermediate Update - Question

Consider the various transactions in the case:

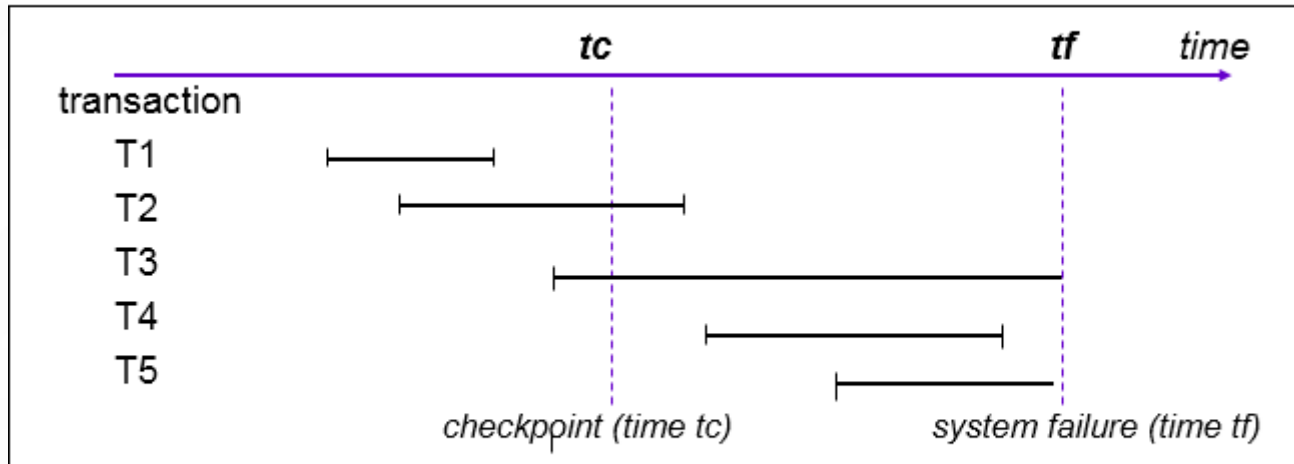


**Question:**

Which transaction(s) need to be undone, redone, or done nothing?

# Intermediate Update - Question

Answer:



**REDO:** T2, T4

**Do nothing:** T1

**UNDO:** T3, T5

# Shadow Paging

---

**No UNDO**  
**No REDO**  
**approach to**  
**RECOVERY**

# Shadow Paging - Idea

---

- ❑ Database is partitioned into fixed-length blocks referred to as PAGES.
- ❑ Page table contains n entries
- ❑ Maintain 2 pages tables during life of transaction
  - The current page table
  - The shadow page table

# Shadow Paging - Process

---

- ❑ When transaction starts BOTH page tables are IDENTICAL
- ❑ The SHADOW page is NEVER changed over the duration of transaction
- ❑ The CURRENT page table may be changed when transaction performs WRITE operation

# Shadow Paging - Example

Transaction T

read x  
write x  
read y  
read z  
write z

Shadow Page Table (links to original values)

x	—
y	—
z	—

Current Page Table (links to updated values)

x	—
y	—
z	—

Disk block


.....
x (original)
.....
y (original)
.....
z (original)
.....
.....
z (updated)
.....
x (updated)
.....



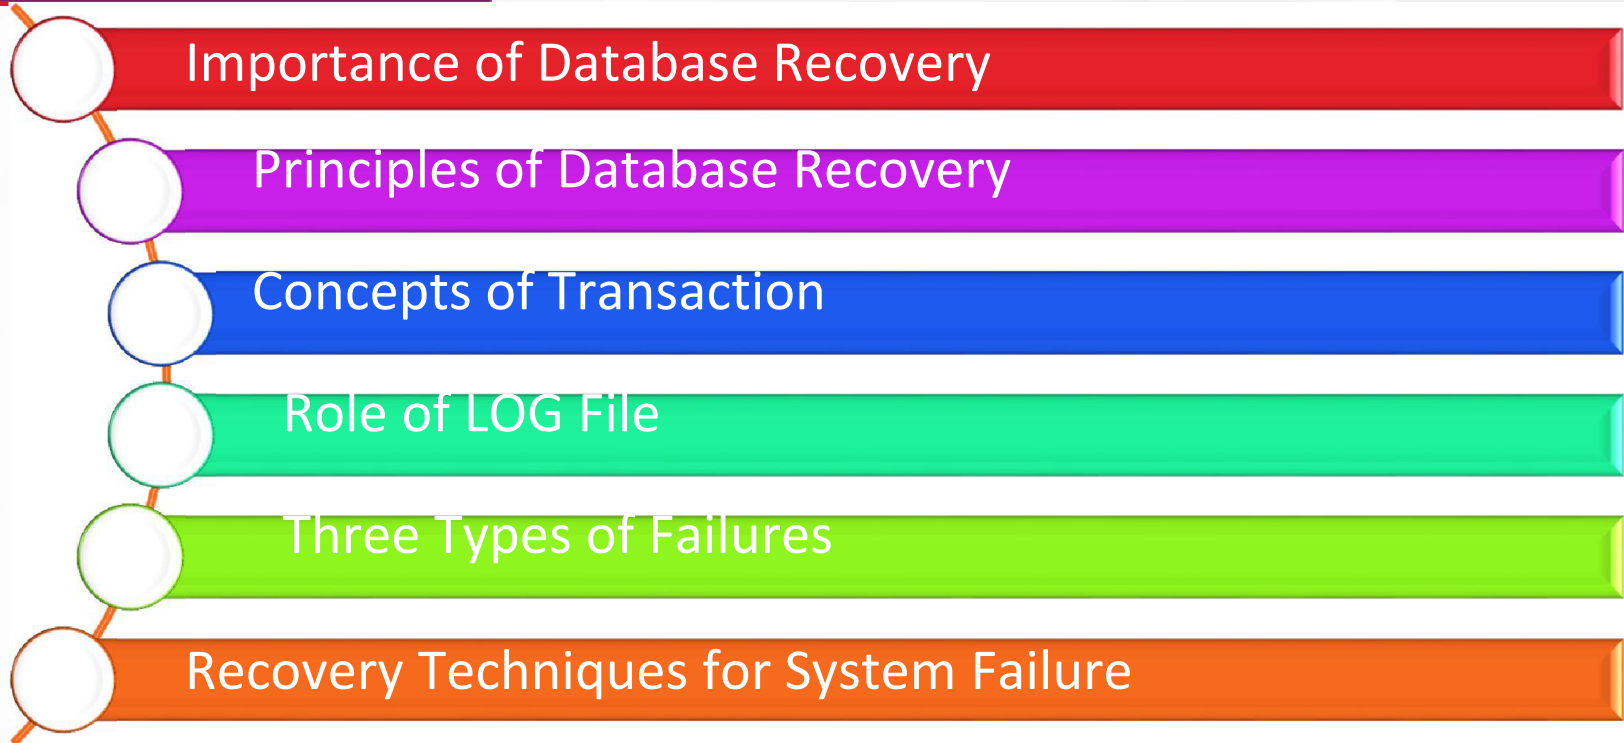
# Shadow Paging - Example

**At commit:-** delete shadow page table  
current page table becomes new 'shadow' page table

**At failure:-** delete current page table  
shadow page table provides the original data values



# SUMMARY



# Thank You

